

Healthcare AI Analytics - Evidence Record

Purpose

Answer supported healthcare business questions from prepared, validated assets while refusing unsupported interpretations.

Work documented

Profiling and cleaning feed semantic definitions, control-total checks and curated knowledge assets. Governed Q&A classifies answerability before producing charts and reviewable reports.

Documented examples and outcomes

The saved case study describes a supported facility case-mix-index answer with benchmark context, chart controls and a supporting table. It also records an unsupported mortality question being refused because required outcome fields were absent, accompanied by a research plan rather than an invented answer. The source reports passing project tests, offline guardrail checks, row-count reconciliation and duplicate-identifier checks. No raw execution log or fresh test run is included in this evidence package.

Supporting artifact

Architecture overview (supporting image available in the website evidence section) - an existing explanatory diagram, not a live application screenshot.

What this supports

A documented readiness-before-AI approach and explicit separation of answerable questions from unsupported claims.

Source basis

Professional Portfolio Platform, "Agentic AI Healthcare Analytics Platform" public-safe case study, dated June 29, 2026; existing healthcare architecture diagram.

Limits

Graduate prototype, not a production clinical, diagnosis, reimbursement or financial decision system. This record does not establish clinical accuracy or savings. Raw healthcare data and private course files are not published.